

PSO2PX USER GUIDE - v1.0.0

PSO2PX Tool

User Guide

“The PSO to Portworx Volume Migration Automation Tool”

Contents

Introduction	3
Prerequisites	3
Limitation and Scope	3
How to run the pso2px tool	7
Usage - migration	7
Usage - migration start (sub-command)	7
Usage - migration resume (sub-command)	8
Usage - migration rollback (sub-command)	8
Usage - version	9
Usage - help	9
Logging	10
Migration Workflow	10
Uninstall PSO	11
Best Practices	11
Known Issues	12



Introduction

The *pso2px* tool is a containerized, python automation tool. *pso2px* is designed to help PSO customers migrate PSO volumes to Portworx. It converts PSO's Kubernetes storage objects, including StorageClasses, PVs and PVCs, into Portworx. This allows the Portworx driver to take over the driver control to continue volume CURD operations and data IOs. *pso2px* requires Portworx version 2.11 or greater, which includes *FlashArray DirectAccess (FADA)* and *FlashBlade DirectAccess (FBDA)* features.

pso2px supports *pure-block* and *pure-file* type volumes migration. It can migrate any typical Kubernetes application types, such as:

- pod
- daemonset
- deployment
- statefulset

pso2px only converts Kubernetes objects, and does not copy data during migration. Usually, migrating one k8s namespace only requires one or two minutes. Before starting, *pso2px performs* a preflight check to ensure that your PSO and Portworx setups are qualified for migration.

Prerequisites

- PSO 6.x.x installed
- Portworx Operator 1.9.0+ installed
- Portworx 2.11.2+ installed
- Applications actively use PSO PVC's to be scaled down manually.

Limitation and Scope

- Application downtime: The migration tool migrates the PSO driver to the Portworx driver (control plane), and results in application downtime due to Kubernetes constraints. Prior to migrating a namespace, you must scale down **all** applications in that specific namespace, run the migration tool, and scale up the application.
- *pso2px* does not support Snapshot object migration.
- PSO 5 migration is not supported. We plan to support it in the v1.1.0 version.
- *pso2px* cannot migrate any volume sizes greater than 40TiB. We plan to support it in the v1.1.0 version.

Tool Installation

1. Save the following yaml spec to *pso2px.yaml* file.

```
wget https://raw.githubusercontent.com/purestorage/pso-csi/master/pso2px/v1.0.0/pso2px.yaml
```



```

# PSO to PX Tool Spec
apiVersion: v1
kind: ServiceAccount
metadata:
  name: pso2px-tool
---
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: pso2px-tool
rules:
  - verbs:
      - '*'
    apiGroups:
      - core.libopenstorage.org
    resources:
      - '*'
  - verbs:
      - get
      - list
      - create
      - update
      - delete
    apiGroups:
      - storage.k8s.io
    resources:
      - storageclasses
      - csinodes
      - csidrivers
  - verbs:
      - get
      - list
      - create
      - delete
      - update
      - patch
    apiGroups:
      - ''
    resources:
      - persistentvolumeclaims

```



```

    - persistentvolumes
    - namespaces
- verbs:
    - get
    - list
  apiGroups:
    - ''
  resources:
    - secrets
    - pods
- verbs:
    - create
  apiGroups:
    - ''
  resources:
    - pods/exec
- verbs:
    - get
    - list
  apiGroups:
    - apps
  resources:
    - deployments
    - statefulsets
    - replicaset
    - daemonsets
- verbs:
    - get
    - list
  apiGroups:
    - batch
  resources:
    - jobs
    - cronjobs
---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: pso2px-tool
subjects:
```



```

- kind: ServiceAccount
  name: pso2px-tool
  namespace: default
roleRef:
  kind: ClusterRole
  name: pso2px-tool
  apiGroup: rbac.authorization.k8s.io
---
apiVersion: v1
kind: Pod
metadata:
  name: pso2px
  labels:
    app: pso2px
spec:
  containers:
  - name: pso2px
    image: portworx/pso2px:v1.0.0
    imagePullPolicy: Always
    volumeMounts:
    - mountPath: /csi.sock
      name: csi-volume
    - mountPath: /usr/local/pso2px/k8s_backups
      name: backup-path
  volumes:
  - name: csi-volume
    hostPath:
      path: /var/lib/kubelet/plugins/pxd.portworx.com/csi.sock
  - name: backup-path
    hostPath:
      path: /usr/local/pso2px/k8s_backups
      type: DirectoryOrCreate
  serviceAccount: pso2px-tool
  serviceAccountName: pso2px-tool

```

2. Apply the yaml file to create a running pod of the tool:

```
kubect! apply -f pso2px.yaml
```

3. check if the pod is up and running:



```
kubect! describe pod/pso2px
```

How to run the pso2px tool

From your terminal, run the following command to login to the container's shell environment:

```
kubect! exec -it pso2px -- bash
```

Usage - migration

The pso2px tool's migration function has 3 major sub-commands: `start`, `rollback`, and `resume`.

```
root@pso2px:/# pso2px migration --help
Usage: pso2px migration [OPTIONS] COMMAND [ARGS]...

Perform migration of PSO Objects to Portworx

Options:
  --help  Show this message and exit.

Commands:
  resume    Resumes migrating PSO objects to Portworx
  rollback  Rollback the Portworx objects back to PSO
  start     Start migration of PSO objects to Portworx
root@pso2px:/#
```

Usage - migration start (sub-command)

```
root@pso2px:/# pso2px migration start --help
Usage: pso2px migration start [OPTIONS]

Start migration of PSO objects to Portworx

Options:
  -n, --namespace TEXT  kubernetes namespace in the cluster
  --help                Show this message and exit.
root@pso2px:/#
```

pso2px can migrate PSO objects either in a particular Kubernetes namespace(s), or all the available namespaces.

To perform migration across all namespaces at once, enter the following command:

```
pso2px migration start
```

Pure Storage does not recommend migrating all namespaces at once.



Instead, you can migrate PSO objects by specifying one or more namespaces explicitly. Migrating namespace-by-namespace can limit the impact of your service downtime only with the namespaces explicitly specified. The services in other namespaces will still be up and running while migrating the current namespace(s).

To perform migration on a particular namespace(s), you must pass on the namespace(s) as an option to the migration command:

```
pso2px migration start --namespace <k8s-namespace>
```

```
pso2px migration start --namespace <k8s-namespace-1> --namespace <k8s-namespace-2>
```

Note: Please make sure that you scale down all applications in the namespace(s) before you start the migration. Tool has the preflight check to check it for you. The preflight check will stop if it detects your application's replica is greater than 0.

Usage - migration resume (sub-command)

```
root@pso2px:/# pso2px migration resume --help
Usage: pso2px migration resume [OPTIONS]

Resumes migrating PSO objects to Portworx

Options:
  --help  Show this message and exit.
root@pso2px:/#
```

During migration, if for any reason the migration is interrupted, then the pso2px tool will stop the migration. The interrupt could be either manually triggered or could be due to any infrastructure related issues.

To continue the migration from the point where it left off, enter the following command:

```
pso2px migration resume
```

This command doesn't accept any options other than `--help`. This resume command will continue migration on namespace(s) which were attempted earlier.

Usage - migration rollback (sub-command)




```

root@pso2px:/# pso2px migration rollback --help
Usage: pso2px migration rollback [OPTIONS]

Rollback the Portworx objects back to PSO

Options:
  -n, --namespace TEXT  kubernetes namespace in the cluster
  --help                Show this message and exit.
root@pso2px:/# █

```

Because of some unforeseen reasons, if the tool fails to migrate the PSO objects to Portworx, then the partially migrated Portworx objects can be rolled back to PSO and applications can then be scaled up to continue with PSO. The rollback helps to minimize the application down time in such situations.

The rollback can also be used on namespaces which successfully migrated from PSO to Portworx. Doing so, will convert all the migrated Portworx objects back to PSO. Pure Storage does not recommend rolling back when the migration is successful.

To perform rollback on a particular namespace, you must pass on the namespace as an option to the rollback sub-command:

```
pso2px migration rollback --namespace <k8s-namespace>
```

You can also rollback multiple namespaces at once (see Best Practices):

```
pso2px migration rollback --namespace <k8s-namespace-1> --namespace <k8s-namespace-2>
```

Usage - version

Once you are in the shell, you can now launch the pso2px migration tool.

To see which the version of pso2px you're running, enter the following command:

```
pso2px --version
```

```

root@pso2px:/# pso2px --version
pso2px : v1.0.0
root@pso2px:/# █

```

Usage - help

To view the pso2px usage options, enter the following command:

```
pso2px --help
```



```

root@pso2px:/# pso2px --help
Usage: pso2px [OPTIONS] COMMAND [ARGS]...

Options:
  -v, --version  Show the version and exit
  --help        Show this message and exit.

Commands:
  migration  Perform migration of PSO Objects to Portworx
root@pso2px:/#

```

Logging

pso2px prints the progress of the tool to the console, and doesn't redirect anything to a log file. To redirect the output to both the console and to a log file, you can pipe the pso2px output to the tee command:

```
pso2px migration start -n <namespace> 2>&1 | tee pso2px.log
```

Migration Workflow

We explain the details of the migration automation process in the tool for people to understand what they would expect during and after migration. The migration has four stages:

1. Preflight Check
2. Application Status Check
3. Backup PSO Objects
4. Migrate PSO Objects to Portworx

Preflight Check:

During the preflight check stage, the tool checks for the prerequisites in the cluster. Your cluster should be running PSO 6.0+ and it should be installed with Portworx 2.11.2+. The Portworx installation should be the operator-based installation. If any of these prerequisites are not satisfied, the tool exits giving the appropriate reason.

Application Status Check:

The pso2px tool expects the application pods, which are actively using the PSO PVC's to be scaled down manually. As the underlying storage driver of the StorageClass object is getting changed during the migration, the applications which are actively using those PVC's need to be scaled down. This is how Kubernetes expects during the storage driver migration.

If the migration is initiated for a particular namespace, then the application pods (Single pod applications, Deployments, StatefulSets, DaemonSets, Jobs, CronJobs) actively using the PSO PVC's in the same namespace, needs to be scaled down.



If the applications are not scaled down, then the pso2px tool, when run, will present a summary of the applications which are actively using the PSO PVCs.

Backup PSO Objects:

After ensuring that no application pods are actively using the PVC's in the namespace where the migration is initiated, the tool will backup the existing PSO Storage Class, PVCs and PVs. These backups will serve as an option to recreate the PSO objects if the migration runs into any issues.

Migrate PSO Objects to Portworx:

After successful backup of existing PSO objects, the tool starts the migration for the given namespace. When the migration starts,

- Deletes the PSO StorageClass
- Creates the Portworx StorageClass using the PSO StorageClass as its source. Except for some StorageClass attributes like the provisioner, backend, other attributes are retained for the Portworx StorageClass.
- Before deleting the PVC, the tool checks the volume reclaim policy of the PVC. If the policy is set to Delete, then the tool changes it to "Retain". The idea here is not to delete the backend volume during migration. If the policy is already set to Retain, it is retained as is.
- After updating the volume reclaim policy, the PSO PVC is deleted.
- Following that, the PSO PV is also deleted.
- Portworx PV is created using PSO PV as its source. Except for some PV attributes like spec→csi→driver and spec→csi→volumeAttributes→backend, other attributes are retained for Portworx PV.
- Portworx PVC is created using the PSO PVC as its source.

This process is repeated for all the PSO PVC & PV in the given namespace. When no namespace is given, this process gets repeated for all the PSO PVC & PV across all the namespaces, one by one.

Uninstall PSO

After migration, you can uninstall PSO by giving `helm uninstall pure-pso`. When uninstalled, the PX migrated default StorageClass (pure-block & pure-file) also gets removed when uninstalling PSO

- When the migrated PX StorageClasses gets deleted, it can be recreated with the storage class definitions made available by the pso2px tool at `/usr/local/pso2px/k8s_backups/migrated_storage_class` directory. This directory will have the JSON definitions of the successfully migrated Storage Class objects.
- To recreate, run this command :
 - `kubectl apply -f /usr/local/pso2px/k8s_backups/migrated_storage_class/pure-file.json`
 - `kubectl apply -f /usr/local/pso2px/k8s_backups/migrated_storage_class/pure-block.json`

Best Practices

- Migrate namespace by namespace. Doing so, the applications running on other namespaces can continue to be online while migration is in progress in the chosen namespace
- When migrating applications by specifying namespace(s), try to have all the applications using the same Pure storage class. If a Pure StorageClass is used by PVC's across more than one



namespace, then you should migrate the applications by giving all of these namespace(s) at once. Refer the section **Usage - migration start (sub-command)** for the command usage

- Similarly, if there is a situation to do a rollback, it is always a best practice to provide the same list of namespace(s) for which the migration was attempted. Refer the section **Usage - migration rollback (sub-command)** for the command usage
- Whenever the PX StorageCluster pods are restarted, the pso2px pod should also be restarted. The pso2px pod uses the CSI Socket file descriptors created by the PX StorageCluster pods. When the file descriptor changes, the pso2px pod should be restarted to pick up the new file descriptor. If we don't do this, migration will fail with errors related to socket availability.

Known Issues

- NEWSTACK-530 : After migration, when you do a rollback to PSO, the PVC & PV gets converted back to PSO, but the portworx cli (pxctl) continues to display the volume when it is listed using `pxctl volume list`
- PWX-25792 : Portworx doesn't support `nosuid` mount_option. Any PSO volume with `nosuid` mount_option will fail to mount after migration
 - To mitigate, before migration, remove `nosuid` mount_option from the storage class definitions
- NEWSTACK-415 : Rawblock volume mode is not supported in FlashArray Direct Access. You should not migrate a PSO rawblock volume to Portworx.

