

PSO2PX USER GUIDE - v1.1.0

PSO2PX Tool

User Guide

“The PSO to Portworx Volume Migration Automation Tool”

Contents

Introduction	3
Prerequisites	3
Limitation and Scope	3
Migrate using pso2px	7
Prerequisites	7
Procedure	7
Other Operations	8
Resume Migration	8
Rollback Migration	9
Version	9
Help	10
Logging	10
Migration Workflow	10
Uninstall PSO	11
Recreate Default StorageClass	11
Best Practices	12
Known Issues	12



Introduction

The *pso2px* tool is a containerized, python automation tool. *pso2px* is designed to help PSO customers in migrating PSO and Flex volumes to Portworx. It converts PSO's Kubernetes storage objects, including StorageClasses, PVs and PVCs, into Portworx. This allows the Portworx driver to take over the driver control to continue volume CURD operations and data IOs. *pso2px* requires Portworx version 2.11 or greater, which includes *FlashArray DirectAccess (FADA)* and *FlashBlade DirectAccess (FBDA)* features.

pso2px supports *pure-block* and *pure-file* type volumes migration. It can migrate any typical Kubernetes application types, such as:

- pod
- daemonset
- deployment
- statefulset

pso2px only converts Kubernetes objects, and does not copy data during migration. Usually, migrating one k8s namespace only requires one or two minutes. Before starting, *pso2px* performs a preflight check to ensure that your PSO and Portworx setups are qualified for migration.

Supported Kubernetes Versions:

- K8S 1.18, 1.20, 1.21 & 1.22

Prerequisites

- PSO 6.x.x or PSO 5.x.x or FlexVolume Plugin installed
- Portworx Operator 1.9.0+ installed
- Portworx 2.11.2+ installed
- Applications actively use PSO/Flex PVC's to be scaled down manually.

Note:

- If you do not have Portworx installed, please reach out to the support team to guide you through the installation procedure.
- We use PSO/Flex terms interchangeably in the rest of this document.

Limitation and Scope

- Application downtime: The migration tool migrates the PSO/Flex driver to the Portworx driver (control plane), and results in application downtime due to Kubernetes constraints. Prior to migrating a namespace, you must scale down **all** applications in that specific namespace, run the migration tool, and scale up the application.



- pso2px does not support Snapshot object migration.
- pso2px cannot migrate any volume sizes greater than 40TiB. We plan to support it in the later versions of the tool.

Tool Installation

1. Save the following yaml spec to pso2px.yaml file.

If you are running k8s version 1.18 or above

wget <https://raw.githubusercontent.com/purestorage/pso-csi/master/pso2px/v1.1.0/pso2px.yaml>

if you are running k8s version 1.17 or below

wget https://raw.githubusercontent.com/purestorage/pso-csi/master/pso2px/v1.1.0/pso2px_k8s_1.17.yaml

```
# PSO to PX Tool Spec
apiVersion: v1
kind: ServiceAccount
metadata:
  name: pso2px-tool
---
kind: ClusterRole
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: pso2px-tool
rules:
- verbs:
  - '*'
  apiGroups:
  - core.libopenstorage.org
  resources:
  - '*'
- verbs:
  - get
  - list
  - create
  - update
  - delete
  apiGroups:
  - storage.k8s.io
  resources:
  - storageclasses
  - csinodes
  - csidrivers
```



```
- verbs:
  - get
  - list
  - create
  - delete
  - update
  - patch
apiGroups:
  - ''

resources:
  - persistentvolumeclaims
  - persistentvolumes
  - namespaces
- verbs:
  - get
  - list
apiGroups:
  - ''

resources:
  - secrets
  - pods
- verbs:
  - create
apiGroups:
  - ''

resources:
  - pods/exec
- verbs:
  - get
  - list
apiGroups:
  - apps
resources:
  - deployments
  - statefulsets
  - replicaset
  - daemonsets
- verbs:
  - get
  - list
```



```

    apiGroups:
      - batch
    resources:
      - jobs
      - cronjobs
---
kind: ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: pso2px-tool
subjects:
- kind: ServiceAccount
  name: pso2px-tool
  namespace: default
roleRef:
  kind: ClusterRole
  name: pso2px-tool
  apiGroup: rbac.authorization.k8s.io
---
apiVersion: v1
kind: Pod
metadata:
  name: pso2px
  labels:
    app: pso2px
spec:
  containers:
  - name: pso2px
    image: portworx/pso2px:v1.1.0
    imagePullPolicy: Always
    volumeMounts:
      - mountPath: /csi.sock
        name: csi-volume
      - mountPath: /usr/local/pso2px/k8s_backups
        name: backup-path
      - mountPath: /usr/local/pso2px/logs
        name: log-path
  volumes:
  - name: csi-volume
    hostPath:

```



```

    path: /var/lib/kubelet/plugins/pxd.portworx.com/csi.sock
  - name: backup-path
    hostPath:
      path: /usr/local/pso2px/k8s_backups
      type: DirectoryOrCreate
  - name: log-path
    hostPath:
      path: /usr/local/pso2px/logs
      type: DirectoryOrCreate
  serviceAccount: pso2px-tool
  serviceAccountName: pso2px-tool

```

2. Apply the yaml file to create a running pod of the tool:

```
kubectl apply -f pso2px.yaml
```

3. check if the pod is up and running:

```
kubectl describe pod/pso2px
```

Migrate using pso2px

Prerequisites

Ensure to scale down all applications in the namespace(s) before you start migration. The tool runs a preflight check, which will stop if it detects your application's replica is greater than 0.

Procedure

From your terminal, run the following command to login to the container's shell environment:

```
kubectl exec -it pso2px -- bash
```

The pso2px tool's migration function has 3 major sub-commands: `start`, `rollback`, and `resume`.



```

root@pso2px:/# pso2px migration --help
Usage: pso2px migration [OPTIONS] COMMAND [ARGS]...

Perform migration of PSO Objects to Portworx

Options:
  --help Show this message and exit.

Commands:
  resume    Resumes migrating PSO objects to Portworx
  rollback  Rollback the Portworx objects back to PSO
  start     Start migration of PSO objects to Portworx
root@pso2px:/#

```

Migrate PSO objects either in a particular kubernetes namespace(s), or all the available namespaces.

```

root@pso2px:/# pso2px migration start --help
Usage: pso2px migration start [OPTIONS]

Start migration of PSO objects to Portworx

Options:
  -n, --namespace TEXT kubernetes namespace in the cluster
  --help                Show this message and exit.
root@pso2px:/#

```

Pure Storage recommends migrating the PSO objects by specifying one or more namespaces explicitly. Migrating namespace-by-namespace can limit the impact of your service downtime only with the namespaces explicitly specified. The services in other namespaces will still be up and running while migrating the current namespace(s).

To perform migration on a particular namespace(s), you must pass the namespace(s) as an option to the migration command:

```
pso2px migration start --namespace <k8s-namespace>
```

```
pso2px migration start --namespace <k8s-namespace-1> --namespace <k8s-namespace-2>
```

Other Operations

Resume Migration




```

root@pso2px:/# pso2px migration resume --help
Usage: pso2px migration resume [OPTIONS]

Resumes migrating PSO objects to Portworx

Options:
  --help  Show this message and exit.
root@pso2px:/# █

```

During migration, if for any reason the migration is interrupted, then the pso2px tool will stop the migration. The interruption could be either manually triggered or due to any infrastructure related issues.

To continue the migration from the point where it left off, enter the following command:

```
pso2px migration resume
```

This command doesn't accept any options other than `-help`. This resume command will continue migration on namespace(s) which were attempted earlier.

Rollback Migration

```

root@pso2px:/# pso2px migration rollback --help
Usage: pso2px migration rollback [OPTIONS]

Rollback the Portworx objects back to PSO

Options:
  -n, --namespace TEXT  kubernetes namespace in the cluster
  --help                Show this message and exit.
root@pso2px:/# █

```

Because of some unforeseen reasons, if the tool fails to migrate the PSO objects to Portworx, then the partially migrated Portworx objects can be rolled back to PSO and applications can then be scaled up to continue with PSO. The rollback helps to minimize the application down time in such situations.

To perform rollback on a particular namespace, you must pass on the namespace as an option to the rollback sub-command:

```
pso2px migration rollback --namespace <k8s-namespace>
```

You can also rollback multiple namespaces at once (see Best Practices):

```
pso2px migration rollback --namespace <k8s-namespace-1> --namespace <k8s-namespace-2>
```

Version

To see which the version of pso2px you're running, enter the following command:



```
pso2px --version
```

```
root@pso2px:/# pso2px --version
pso2px : v1.1.0
root@pso2px:/#
```

Help

To view the pso2px usage options, enter the following command:

```
pso2px --help
```

```
root@pso2px:/# pso2px --help
Usage: pso2px [OPTIONS] COMMAND [ARGS]...

Options:
  -v, --version  Show the version and exit
  --help         Show this message and exit.

Commands:
  migration  Perform migration of PSO Objects to Portworx
root@pso2px:/#
```

Logging

pso2px prints the progress of the tool to the console, and also writes them to a log file. The log files are available at `/usr/local/pso2px/logs` directory.

Migration Workflow

Following is the migration workflow. It has four stages:

1. Preflight Check
2. Application Status Check
3. Backup PSO Objects
4. Migrate PSO Objects to Portworx

Preflight Check:

During the preflight check stage, the tool checks for the prerequisites in the cluster. Your cluster must be running PSO or FlexVolume Plugin and it should be installed with Portworx **version 2.11.2 or newer**. The Portworx installation should be the operator-based installation. If any of these prerequisites are not satisfied, the tool **will** exit with **an** appropriate reason.

Application Status Check:



The pso2px tool expects the application pods, which are actively using the PSO PVCs, to be scaled down manually. As the underlying storage driver of the StorageClass object is getting changed during the migration, the applications which are actively using these PVCs need to be scaled down. **This is a Kubernetes requirement during the storage driver migration.**

If the migration is initiated for a particular namespace, then the application pods (Single pod applications, Deployments, StatefulSets, DaemonSets, Jobs, CronJobs) actively using the PSO PVCs in the same namespace, needs to be scaled down.

If the applications are not scaled down, then the pso2px tool, when run, will present a summary of the applications which are actively using the PSO PVCs.

Backup PSO Objects:

After ensuring that no application pods are actively using the PVC's in the namespace where the migration is initiated, the tool will backup the existing PSO StorageClass, PVCs and PVs. These backups will serve as an option to recreate the PSO objects if the migration is unsuccessful.

Migrate PSO Objects to Portworx:

After a successful backup of the existing PSO objects, the tool starts the migration for the given namespace. When the migration starts, the tool:

- Deletes the PSO StorageClass
- Creates the Portworx StorageClass using the PSO StorageClass as its source. Except for some StorageClass attributes like the provisioner, backend, other attributes are retained for the Portworx StorageClass.
- Verifies the volume reclaim policy of the PVC. If the policy is set to Delete, then the tool changes it to "Retain". The idea here is not to delete the backend volume during migration. If the policy is already set to Retain, it is retained as is.
- After updating the volume reclaim policy, the PSO PVC is deleted.
- Following that, deletes the PSO PV.
- Creates Portworx PV using PSO PV as its source. Except for some PV attributes like spec→csi→driver and spec→csi→volumeAttributes→backend, other attributes are retained for Portworx PV.
- Creates Portworx PVC using PSO PVC as its source.

This process is repeated for all the PSO PVC and PV in a given namespace. When no namespace is given, this process gets repeated for all the PSO PVC and PV across all the namespaces, one by one.

Uninstall PSO

Run the following command to uninstall PSO.

For uninstalling PSO6, run

```
helm uninstall pure-pso
```

For uninstalling PSO5, run

```
helm uninstall pure-csi
```

For uninstalling Flex VolumePlugin, run



```
helm uninstall pure-storage-driver
```

When you uninstall PSO, the PX migrated default StorageClass (pure-block & pure-file) also gets removed.

Recreate Default StorageClass

When the migrated PX StorageClasses gets deleted, it can be recreated with the storage class definitions made available by the pso2px tool at `/usr/local/pso2px/k8s_backups/migrated_storage_class` directory. This directory will have the JSON definitions of the successfully migrated Storage Class objects.

To recreate, run the following command :

- `kubectl apply -f /usr/local/pso2px/k8s_backups/migrated_storage_class/pure-file.json`
- `kubectl apply -f /usr/local/pso2px/k8s_backups/migrated_storage_class/pure-block.json`

Best Practices

- Migrate namespace by namespace. Doing so, the applications running on other namespaces can continue to be online while migration is in progress in the chosen namespace
- When migrating applications by specifying the namespace(s), try to have all the applications using the same Pure StorageClass. When a Pure StorageClass is used by PVC's across multiple namespaces, use the following command to migrate these namespaces all at once: `pso2px migration start --namespace <k8s-namespace-1> --namespace <k8s-namespace-2>`
- Similarly, if there is a situation to do a rollback, it is always a best practice to provide the same list of namespace(s) for which the migration was attempted. Like, `pso2px migration rollback --namespace <k8s-namespace-1> --namespace <k8s-namespace-2>`
- Whenever the PX StorageCluster pods are restarted, the pso2px pod should also be restarted. The pso2px pod uses the CSI Socket file descriptors created by the PX StorageCluster pods. When the file descriptor changes, the pso2px pod should be restarted to pick up the new file descriptor. Otherwise, migration will fail with errors related to socket availability.

Known Issues

- NEWSTACK-530 : After migration, when you do a rollback to PSO, the PVC & PV gets converted back to PSO, but the Portworx CLI utility (pxctl) continues to display the volume when it is listed using `pxctl volume list`
- PWX-25792 : Portworx doesn't support `nosuid` mount_option. Any PSO volume with `nosuid` mount_option will fail to mount after migration
 - To migrate, remove the `nosuid` mount_option from the StorageClass definition before initiating migration process
- PWX-24890 : After migrating a Flashblade Direct Access volume to Portworx, and scaling up the associated application, `pxctl volume list` still shows the volume as detached. But the volume will be mounted to the application pod and allows IO operations
- NEWSTACK-415 : Rawblock volume mode is not supported in FlashArray Direct Access. You should not migrate a PSO rawblock volume to Portworx.

